



# MatchIT: Python for Data Science

## Lecture: Pandas

# Lecture 2: Content

---

- Data Manipulation with Pandas
- Pandas Objects
- Indexing and Data Selection on Pandas Objects
- Operating on Data
  - Universal Functions and Index Alignment
  - Handling Missing Data
  - Aggregation
  - Groupby
    - Filter
    - Transform
    - Apply
  - Pivot Tables

# Data Manipulation with Pandas

---

- Pandas is a newer package built on NumPy
- Why Pandas?
  - NumPy's ndarray data structure contains essential features needed to store and manipulate data in numerical computing tasks - BUT
  - Numpy has some limitations that Pandas tries to solve, such as
    - Greater flexibility to attach labels to data
    - Simplifies task of working with missing data
    - Operations such as groupings and pivots
- Pandas Series and DataFrame objects are built on NumPy array structure to provide

# Pandas Objects

---

- Pandas objects are an enhanced versions of NumPy structured arrays
  - Rows and columns are identified with labels rather than simple integer indices.
  - Provide useful tools, methods, and functionality on top of the basic data structures
  - Understanding the basic data structures is essential in taking full advantage of Pandas package

# Pandas Data Structures

---

- Three fundamental Pandas Data structures
  - Series
  - DataFrame
  - Index

# Pandas Data Structures: Series Object

---

- Can be created from a list or array

```
from scipy import *
from matplotlib.pyplot import *

import numpy as np
import pandas as pd

data= pd.Series([0.25, 0.5, 0.75, 1.0])
data
```

- The series output, data, combines both a sequence and a sequence of indices.

```
In [2]: data
Out[2]:
0      0.25
1      0.50
2      0.75
3      1.00
dtype: float64
```

# Pandas Data Structures: Series Object

- We can access easily data values:

```
In [3]: data.values  
Out[3]: array([0.25, 0.5 , 0.75, 1.  ])
```

- As well as indices:

```
In [4]: data.index  
Out[4]: RangeIndex(start=0, stop=4, step=1)
```

- As with NumPy array, data can be access via index

```
In [5]: data[1]  
Out[5]: 0.5
```

```
In [6]: data[1:3]  
Out[6]:  
1    0.50  
2    0.75  
dtype: float64
```

- Series object similar to one-dimensional NumPy array, but Series object offers more..

# Pandas Data Structures: Series Object

---

- Series object has explicitly defined index associated with the values:

```
In [10]: data= pd.Series([0.25, 0.5, 0.75, 1.0], index=['a', 'b', 'c', 'd'])
```

```
In [11]: data
```

```
Out[11]: |  
a      0.25  
b      0.50  
c      0.75  
d      1.00  
dtype: float64
```

```
In [12]: data['b']
```

```
Out[12]: 0.5
```

- Series can be thought of as specialized Python dictionary (key, value), but with more efficient implementation



# Pandas Data Structures: Series Object

---

- Series can be created from a Python dictionary

```
In [23]: results_dict={'passed': 12,  
...: 'passed with distinction': 5,  
...: 'failed': 2}
```

```
In [24]: results_ser=pd.Series(results_dict)
```

```
In [25]: results_ser
```

```
Out[25]:  
passed                12  
passed with distinction    5  
failed                 2  
dtype: int64
```

# Pandas Data Structures: Data Frame Object

- DataFrame object can also be considered as generalization of NumPy array or specialization of a Python dictionary
  - DataFrame as a generalized NumPy array
    - Analog to *two-dimensional* array with *flexible row indices* and *flexible column names*

*Two-dimensional array = Ordered Sequence of aligned one dimension columns*  
*DataFrame = Ordered Sequence of aligned Series objects*

*aligned\** - share the same index

# Pandas Data Structures: Data Frame Object

- Example:

```
In [38]: results_ser
```

```
Out[38]:
```

```
passed          12
passed with distinction    5
failed           2
dtype: int64
```

```
In [39]: grades_dict={'passed' : 3, 'passed with distinction' : 4, 'failed': 2}
```

```
In [40]: grading_results=pd.DataFrame({'results': results_ser, 'grades': grades_dict})
```

```
In [41]: grading_results
```

```
Out[41]:
```

	results	grades
failed	2	2
passed	12	3
passed with distinction	5	4

```
In [42]: grading_results.index
```

```
Out[42]: Index(['failed', 'passed', 'passed with distinction'], dtype='object')
```

```
In [43]: grading_results.columns
```

```
Out[43]: Index(['results', 'grades'], dtype='object')
```

# Pandas Data Structures: Data Frame Object

---

- How to construct DataFrame Object
  - From a single Series object
  - From a list of dictionaries
  - From a dictionary of Series objects
  - From a two-dimensional NumPy array
  - From a NumPy structured array

# Pandas Data Structures: Data Frame Object

## Construction Examples

---

- From a single Series object:

```
In [45]: pd.DataFrame(results_ser, columns=['results'])
Out[45]:
```

	results
passed	12
passed with distinction	5
failed	2

- From a list of dictionary

```
In [50]: data = [{'a': i, 'b': 2*i, 'c': i + 2*i}
...:             for i in range(3)]
...: pd.DataFrame(data)
Out[50]:
```

	a	b	c
0	0	0	0
1	1	2	3
2	2	4	6

# Pandas Data Structures: Data Frame Object

## *Construction Examples*

- From a dictionary of Series objects

```
In [57]: grades_dict
Out[57]: {'passed': 3, 'passed with distinction': 4, 'failed': 2}

In [58]: grades_ser=pd.Series(grades_dict)

In [59]: pd.DataFrame({'results':results_ser,'grades':grades_ser})
Out[59]:
```

	results	grades
passed	12	3
passed with distinction	5	4
failed	2	2

- From two-dimensional NumPy array

```
In [61]: pd.DataFrame(np.random.rand(3,2),
...:                  columns=['foo', 'bar'],
...:                  index=['a', 'b', 'c'])
Out[61]:
```

	foo	bar
a	0.640737	0.344832
b	0.257366	0.020849
c	0.687164	0.142364

# Pandas Data Structures: Data Frame Object

## *Construction Examples*

---

- From NumPy structured array:

```
In [68]: A=np.zeros(3, dtype=[('A', 'i8'), ('B','f8')])
```

```
In [69]: A
```

```
Out[69]: array([(0, 0.), (0, 0.), (0, 0.)], dtype=[('A', '<i8'), ('B', '<f8')])
```

```
In [70]: pd.DataFrame(A)
```

```
Out[70]:
```

	A	B
0	0	0.0
1	0	0.0
2	0	0.0

# Pandas Data Structures: Index Object

---

- Series and DataFrame objects contain an explicit index for data access and data modification
  - Index : *Immutable array* or *ordered set (or multiset)*

```
In [76]: X=np.arange(2,16,3)
In [77]: ind=pd.Index(X)
In [78]: ind[1]
Out[78]: 5
In [79]: X
Out[79]: array([ 2,  5,  8, 11, 14])
In [80]:
```



# Pandas Data Structures: Index Object

- Index as immutable array having attributes familiar from NumPy arrays:

```
In [83]: X
Out[83]: array([ 2,  5,  8, 11, 14])
```

```
In [84]: ind[0]
Out[84]: 2
```

```
In [85]: ind[::2]
Out[85]: Int64Index([2, 8, 14], dtype='int64')
```

```
In [86]: print(ind.size, ind.shape, ind.ndim, ind.dtype)
5 (5,) 1 int64
```

- Index objects are immutable- can not be modified like:*

```
In [87]: ind[1]=0
Traceback (most recent call last):
```

```
File "<ipython-input-87-b10b243764e2>", line 1, in <module>
    ind[1]=0
```

```
File "/anaconda3/lib/python3.7/site-packages/pandas/core/indexes/base.py", line 3938, in __setitem__
    raise TypeError("Index does not support mutable operations")
```

```
TypeError: Index does not support mutable operations
```

# Pandas Data Structures: Index Object

---

- Indexes support operations typical of data sets: unions, intersections, differences, and other combinations

```
In [90]: indA=pd.Index([1,3,5,7,9])
```

```
In [91]: indB=pd.Index([2,3,5,7,11])
```

```
In [92]: indA & indB
```

```
Out[92]: Int64Index([3, 5, 7], dtype='int64')
```

```
In [93]: indA | indB
```

```
Out[93]: Int64Index([1, 2, 3, 5, 7, 9, 11], dtype='int64')
```

```
In [94]: indA ^ indB
```

```
Out[94]: Int64Index([1, 2, 9, 11], dtype='int64')
```

# Pandas: Data Indexing and Selection

- Review: Remember methods and tools to access, set and modify values in NumPy arrays:

```
In [27]: A=np.arange(10,40)
```

Indexing

```
In [28]: A
```

```
Out[28]:
```

```
array([10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26,
       27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39])
```

```
In [29]: A=A.reshape(5,6)
```

```
In [30]: A
```

```
Out[30]:
```

```
array([[10, 11, 12, 13, 14, 15],
       [16, 17, 18, 19, 20, 21],
       [22, 23, 24, 25, 26, 27],
       [28, 29, 30, 31, 32, 33],
       [34, 35, 36, 37, 38, 39]])
```

```
In [31]: A=A[:,1:5]
```

Slicing

```
In [32]: A
```

```
Out[32]:
```

```
array([[11, 12, 13, 14],
       [17, 18, 19, 20],
       [23, 24, 25, 26],
       [29, 30, 31, 32],
       [35, 36, 37, 38]])
```

```
In [33]: A=A>23
```

Masking

```
In [34]: A
```

```
Out[34]:
```

```
array([[False, False, False, False],
       [False, False, False, False],
       [False,  True,  True,  True],
       [ True,  True,  True,  True],
       [ True,  True,  True,  True]])
```

# Pandas: Data Indexing and Selection

- Review: Remember methods and tools to access, set and modify values in NumPy arrays:

```
In [51]: y=np.arange(35).reshape(5,7)
```

```
In [52]: y
```

```
Out[52]:
```

```
array([[ 0,  1,  2,  3,  4,  5,  6],
       [ 7,  8,  9, 10, 11, 12, 13],
       [14, 15, 16, 17, 18, 19, 20],
       [21, 22, 23, 24, 25, 26, 27],
       [28, 29, 30, 31, 32, 33, 34]])
```

Fancy indexing

```
In [53]: y[1:5:2, ::3]
```

```
Out[53]:
```

```
array([[ 7, 10, 13],
       [21, 24, 27]])
```

```
In [54]:
```

# Pandas: Data Indexing and Selection

- Data Selection in Series

```
In [63]: data=pd.Series([0.25, 0.5, 0.75, 1.0],  
...: index=['a', 'b', 'c', 'd'])
```

```
In [64]: data
```

```
Out[64]:
```

```
a    0.25
```

```
b    0.50
```

```
c    0.75
```

```
d    1.00
```

```
dtype: float64
```

```
In [65]: data['b']
```

```
Out[65]: 0.5
```

Accessing an element

```
In [66]: 'a' in data
```

```
Out[66]: True
```

Examine key or index

```
In [67]: data.keys()
```

```
Out[67]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
In [68]: list(data.items())
```

```
Out[68]: [('a', 0.25), ('b', 0.5), ('c', 0.75), ('d', 1.0)]
```

```
In [69]: data['e'] = 1.25
```

Adding new value

```
In [70]: data
```

```
Out[70]:
```

```
a    0.25
```

```
b    0.50
```

```
c    0.75
```

```
d    1.00
```

```
e    1.25
```

```
dtype: float64
```

# Pandas: Data Indexing and Selection

- Data Selection in Series
  - Series builds on dictionary-like interface
  - Series provide array-style item selection as NumPy arrays: slicing, masking, fancy indexing...

```
In [72]: data['a':'c']
```

```
Out[72]:
```

```
a    0.25
```

```
b    0.50
```

```
c    0.75
```

```
dtype: float64
```

Slicing by explicit index

```
In [73]: data[0:2]
```

```
Out[73]:
```

```
a    0.25
```

```
b    0.50
```

```
dtype: float64
```

Slicing by implicit integer index

```
In [74]: data[(data > 0.3) & (data < 0.8)]
```

```
Out[74]:
```

```
b    0.50
```

```
c    0.75
```

```
dtype: float64
```

Masking

```
In [75]: data[['a', 'e']]
```

```
Out[75]:
```

```
a    0.25
```

```
e    1.25
```

```
dtype: float64
```

Indexing

# Pandas: Data Indexing and Selection

**loc, iloc, and ix** In [77]: data =pd.Series(['a', 'b', 'c'], index=[1,3,5])

In [78]: data

Out[78]:

```
1    a
3    b
5    c
dtype: object
```

In [79]: data[1]

Out[79]: 'a'

Explicit index

In [80]: data[1:3]

Out[80]:

```
3    b
5    c
dtype: object
```

Implicit integer index

In [81]: data.loc[1:3]

Out[81]:

```
1    a
3    b
dtype: object
```

loc-always references  
explicit index

In [82]: data.iloc[1:3]

Out[82]:

```
3    b
5    c
dtype: object
```

iloc-always references  
implicit index

# Pandas: Data Indexing and Selection

## Data Selection in DataFrame as Dictionary

- Review Dataframe :  
Can be viewed as a two dimensional structured array or dictionary of Series structures sharing the same index.

```
In [84]: area = pd.Series({'California': 423967, 'Texas': 695662,
...:                      'New York': 141297, 'Florida': 170312,
...:                      'Illinois': 149995})
...: pop = pd.Series({'California': 38332521, 'Texas': 26448193,
...:                  'New York': 19651127, 'Florida': 19552860,
...:                  'Illinois': 12882135})
...: data = pd.DataFrame({'area':area, 'pop':pop})
...: data
```

Out[84]:

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

Dataframe as dictionary

```
In [85]: data['area']
```

Out[85]:

California	423967
Texas	695662
New York	141297
Florida	170312
Illinois	149995

Name: area, dtype: int64

```
In [86]: data.area
```

Out[86]:

California	423967
Texas	695662
New York	141297
Florida	170312
Illinois	149995

Name: area, dtype: int64

```
In [87]: data.area is data['area']
```

Out[87]: True

```
In [88]: data.pop is data['pop']
```

Out[88]: False



# Pandas: Data Indexing and Selection

```
In [90]: data
```

```
Out[90]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

Data selection in dataframe

```
In [91]: data.iloc[:3,:2]
```

```
Out[91]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127

```
In [92]: data.iloc[:3,:]
```

```
Out[92]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127

```
In [93]: data.loc[:, 'Illinois', : 'pop']
```

```
Out[93]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

```
In [94]: data.loc[:, 'New York', :]
```

```
Out[94]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127

# Pandas: Data Indexing and Selection

---

```
In [100]: data['density']=data['pop']/data['area']
```

```
In [101]: data
```

```
Out[101]:
```

	area	pop	density
California	423967	38332521	90.413926
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

Dataframe as two-dimensional array

```
In [102]: data.values
```

```
Out[102]:
```

```
array([[4.23967000e+05, 3.83325210e+07, 9.04139261e+01],
       [6.95662000e+05, 2.64481930e+07, 3.80187404e+01],
       [1.41297000e+05, 1.96511270e+07, 1.39076746e+02],
       [1.70312000e+05, 1.95528600e+07, 1.14806121e+02],
       [1.49995000e+05, 1.28821350e+07, 8.58837628e+01]])
```

```
In [103]: data.T
```

```
Out[103]:
```

	California	Texas	New York	Florida	Illinois
area	4.239670e+05	6.956620e+05	1.412970e+05	1.703120e+05	1.499950e+05
pop	3.833252e+07	2.644819e+07	1.965113e+07	1.955286e+07	1.288214e+07
density	9.041393e+01	3.801874e+01	1.390767e+02	1.148061e+02	8.588376e+01

```
In [104]: data.values[0]
```

```
Out[104]: array([4.23967000e+05, 3.83325210e+07, 9.04139261e+01])
```

```
In [105]: data['area']
```

```
Out[105]:
```

California	423967
Texas	695662
New York	141297
Florida	170312
Illinois	149995

```
Name: area, dtype: int64
```

## Pandas: Data Indexing and Selection

```
In [107]: data
```

```
Out[107]:
```

	area	pop	density
California	423967	38332521	90.413926
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

Dataframe as two-dimensional array

```
In [108]: data.iloc[:3,:2]
```

```
Out[108]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127

```
In [109]: data.loc[:'Illinois',:'pop']
```

```
Out[109]:
```

	area	pop
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

```
In [110]: data.loc[data.density>100, ['pop', 'density']]
```

```
Out[110]:
```

	pop	density
New York	19651127	139.076746
Florida	19552860	114.806121

# Pandas: Operating on Data in Pandas

- Panda inherits from NumPy functionality that performs efficient element-wise operations for example addition, subtraction, multiplication as well as trigonometric functions, exponential and logarithmic functions, etc.
- Universal Functions, which are used for computation on NumPy Arrays is essential for performing operations. Panda automatically aligns indexes when passing the objects to the functions

```
In [97]: rng=np.random.RandomState(42)
```

```
In [98]: ser=pd.Series(rng.randint(0,10,3))
```

```
In [99]: ser
```

```
Out[99]:  
0    6  
1    3  
2    7  
dtype: int64
```

```
In [100]: df=pd.DataFrame(rng.randint(0,10,(3,4)),  
...: columns=['A','B','C','D'])
```

```
In [101]: df
```

```
Out[101]:  
   A  B  C  D  
0  4  6  9  2  
1  6  7  4  3  
2  7  7  2  5
```

```
In [102]: np.exp(ser)
```

```
Out[102]:  
0    403.428793  
1    20.085537  
2   1096.633158  
dtype: float64
```

```
In [103]: np.sin(df * np.pi / 4)
```

```
Out[103]:  
   A          B          C          D  
0  1.224647e-16 -1.000000  7.071068e-01  1.000000  
1 -1.000000e+00 -0.707107  1.224647e-16  0.707107  
2 -7.071068e-01 -0.707107  1.000000e+00 -0.707107
```

# Pandas: Operating on Data in Pandas

## UFuncs and Index Alignment in Series

```
In [107]: area = pd.Series({'Alaska': 1723337, 'Texas': 695662,  
...:                       'California': 423967}, name='area')  
...: population = pd.Series({'California': 38332521, 'Texas': 26448193,  
...:                        'New York': 19651127}, name='population')
```

```
In [108]: population/area
```

```
Out[108]:
```

Alaska	NaN
California	90.413926
New York	NaN
Texas	38.018740

←

```
dtype: float64
```

```
In [109]: area.index
```

```
Out[109]: Index(['Alaska', 'Texas', 'California'], dtype='object')
```

```
In [110]: population.index
```

```
Out[110]: Index(['California', 'Texas', 'New York'], dtype='object')
```

```
In [111]: area.index | population.index
```

```
Out[111]: Index(['Alaska', 'California', 'New York', 'Texas'], dtype='object')
```

```
In [113]: area.divide(population, fill value=0)
```

```
Out[113]:
```

Alaska	inf
California	0.011060
New York	0.000000
Texas	0.026303

←

```
dtype: float64
```

```
In [114]: population.divide(area, fill value=0)
```

```
Out[114]:
```

Alaska	0.000000
California	90.413926
New York	inf
Texas	38.018740

←

```
dtype: float64
```

# Pandas: Operating on Data in Pandas

## UFuncs and Index Alignment in DataFrame

```
In [116]: A = pd.DataFrame(rng.randint(0, 20, (2, 2)),  
...:                      columns=list('AB'))  
...: A
```

```
Out[116]:
```

```
   A  B  
0   6 18  
1  10 10
```

```
In [117]: B = pd.DataFrame(rng.randint(0, 10, (3, 3)),  
...:                      columns=list('BAC'))  
...: B
```

```
Out[117]:
```

```
   B  A  C  
0   7  4  3  
1   7  7  2  
2   5  4  1
```

```
In [118]: A+B
```

```
Out[118]:
```

```
   A    B  C  
0 10.0 25.0 NaN  
1 17.0 17.0 NaN  
2  NaN  NaN NaN
```

```
In [119]: A.add(B, fill_value=0)
```

```
Out[119]:
```

```
   A    B  C  
0 10.0 25.0 3.0  
1 17.0 17.0 2.0  
2  4.0  5.0 1.0
```

### Python OperatorPandas Method(s)

+

add()

-

sub(), subtract()

\*

mul(), multiply()

/

truediv(), div(), divide()

//

floordiv()

%

mod()

\*\*

pow()

# Pandas: Handling Missing Data

- Real world data used in analysis is rarely clean and homogeneous.
- Python provides two ways to represent missing data:
  - 1) Using **None** Object . Not suited for aggregate functions, `sum()` or `min()` across an array: produces an error
  - 2) **NaN** (not a number), is a special floating –point value recognized by all systems that used the standard IEEE floating-point representation

# Pandas: Handling Missing Data with NaN

```
In [130]: vals2 = np.array([1, np.nan, 3, 4])
```

```
In [131]: vals2.dtype
```

```
Out[131]: dtype('float64')
```

```
In [132]: 1 + np.nan
```

```
Out[132]: nan
```

```
In [133]: 0 * np.nan
```

```
Out[133]: nan
```

```
In [134]: vals2.sum(), vals2.min(), vals2.max()
```

```
Out[134]: (nan, nan, nan)
```

```
In [135]: np.nansum(vals2), np.nanmin(vals2), np.nanmax(vals2)
```

```
Out[135]: (8.0, 1.0, 4.0)
```



# Pandas: Handling Missing Data with NaN

## 1) Detecting null values:

```
In [139]: data = pd.Series([1, np.nan, 'hello', None])
```

```
In [140]: data.isnull()
```

```
Out[140]:
```

```
0    False
1     True
2    False
3     True
dtype: bool
```

```
In [141]: data[data.notnull()]
```

```
Out[141]:
```

```
0      1
2    hello
dtype: object
```

```
In [142]: data.dropna()
```

```
Out[142]:
```

```
0      1
2    hello
dtype: object
```

# Pandas: Handling Missing Data with NaN

## 1) Detecting null values:

```
In [139]: data = pd.Series([1, np.nan, 'hello', None])
```

```
In [140]: data.isnull()
```

```
Out[140]:
```

```
0    False
1     True
2    False
3     True
dtype: bool
```

```
In [141]: data[data.notnull()]
```

```
Out[141]:
```

```
0      1
2    hello
dtype: object
```

```
In [142]: data.dropna()
```

```
Out[142]:
```

```
0      1
2    hello
dtype: object
```

# Pandas: Aggregation

```
In [150]: import seaborn as sns
```

```
In [151]: planets=sns.load_dataset('planets')
```

← Data available through  
seaborn package

```
In [152]: planets.shape
```

```
Out[152]: (1035, 6)
```

```
In [153]: planets.head
```

```
Out[153]:
```

```
<bound method NDFrame.head of
0    Radial Velocity    1    269.300000    7.100    77.40    2006
1    Radial Velocity    1    874.774000    2.210    56.95    2008
2    Radial Velocity    1    763.000000    2.600    19.84    2011
```

```
In [155]: planets.dropna().describe()
```

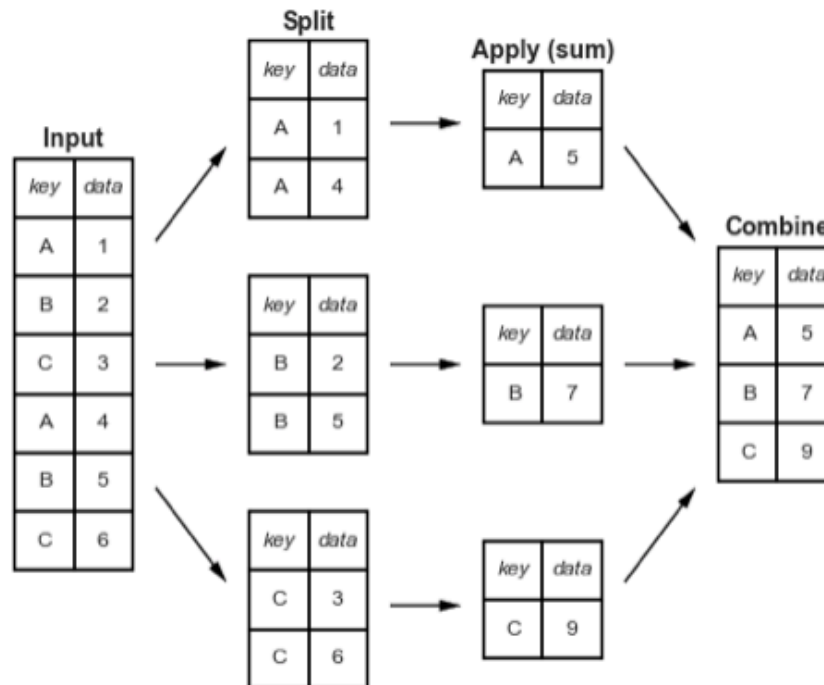
← Understand data with describe()

```
Out[155]:
```

	number	orbital_period	mass	distance	year
count	498.000000	498.000000	498.000000	498.000000	498.000000
mean	1.73494	835.778671	2.509320	52.068213	2007.377510
std	1.17572	1469.128259	3.636274	46.596041	4.167284
min	1.00000	1.328300	0.003600	1.350000	1989.000000
25%	1.00000	38.272250	0.212500	24.497500	2005.000000
50%	1.00000	357.000000	1.245000	39.940000	2009.000000
75%	2.00000	999.600000	2.867500	59.332500	2011.000000
max	6.00000	17337.500000	25.000000	354.000000	2014.000000

# Pandas GroupBy: Split, Apply, Combine

- groupby operation allows us to aggregate on any label or index



# Pandas GroupBy: Split, Apply, Combine

- groupby function on dataframe returns DataFrameGroupBy object.
- to calculate mean, we call mean() function on the DataFrameGroupBy object
- to plot, add call to plot() function

Example:

```
In [165]: df
```

```
Out[165]:
```

	key	data
0	A	0
1	B	1
2	C	2
3	A	3
4	B	4
5	C	5

```
In [166]: groupByKey=df.groupby('key')
```


```
In [167]: groupByKey.mean()
```

```
Out[167]:
```

	data
key	
A	1.5
B	2.5
C	3.5

```
In [168]: groupByKey.mean().plot()
```

```
Out[168]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f61e828>
```



# Pandas GroupBy: Split, Apply, Combine

- Exercise (Stand up and stretch): DataFrame: MatchIT students

Split, Apply, Combine:

	Name	Country	Age	Weight	Height
0	Olof Jonson	Sweden	30	70	170
1	.....				

```
students_df.groupby('Country')
students_df.groupby('Country').shape()
students_df.groupby('Country')['Age'].max()
students_df.groupby('Country')['Age'].min()
```

# Pandas GroupBy: Split, Apply, Combine

```
In [183]: import seaborn as sns
...: planets = sns.load_dataset('planets')
...: planets.shape
...: Out[170]: (1035, 6)
```

```
In [184]: planets.head()
```

```
Out[184]:
```

	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011
3	Radial Velocity	1	326.030	19.40	110.62	2007
4	Radial Velocity	1	516.220	10.50	119.47	2009

```
In [185]: groupByMethod = planets.groupby('method') ← Column Indexing
```

```
In [186]: groupByMethod['year'].median() ←
```

```
Out[186]:
```

method	
Astrometry	2011.5
Eclipse Timing Variations	2010.0
Imaging	2009.0
Microlensing	2010.0
Orbital Brightness Modulation	2011.0
Pulsar Timing	1994.0
Pulsation Timing Variations	2007.0
Radial Velocity	2009.0
Transit	2012.0
Transit Timing Variations	2012.5

GroupByDataFrame  
Object call to median()  
method

```
Name: year, dtype: float64
```

# Pandas GroupBy: Split, Apply, Combine

Iteration over groups:

```
In [246]: for (method, group) in planets.groupby('method'):
...:     print("\n {0:20s} \n {1}".format(method, group.describe()))
...:
```

```
Astrometry
count    number  orbital_period  mass  distance  year
mean      1.0      631.180000   NaN  17.875000  2011.50000
std       0.0      544.217663   NaN   4.094148   2.12132
min       1.0      246.360000   NaN  14.980000  2010.00000
25%       1.0      438.770000   NaN  16.427500  2010.75000
50%       1.0      631.180000   NaN  17.875000  2011.50000
75%       1.0      823.590000   NaN  19.322500  2012.25000
max       1.0     1016.000000   NaN  20.770000  2013.00000

Eclipse Timing Variations
count    number  orbital_period  mass  distance  year
mean    1.666667  4751.644444   5.125000  315.360000  2010.000000
std     0.500000  2499.130945   1.308148  213.203907   1.414214
min     1.000000  1916.250000   4.200000  130.720000  2008.000000
25%     1.000000  2900.000000   4.662500  130.720000  2009.000000
50%     2.000000  4343.500000   5.125000  315.360000  2010.000000
75%     2.000000  5767.000000   5.587500  500.000000  2011.000000
max     2.000000  10220.000000   6.050000  500.000000  2012.000000

Imaging
count    number  orbital_period  mass  distance  year
mean     1.315789 118247.737500   NaN   67.715937  2009.131579
std      0.933035 213978.177277   NaN   53.736817   2.781901
min      1.000000  4639.150000   NaN    7.690000  2004.000000
25%      1.000000  8343.900000   NaN   22.145000  2008.000000
50%      1.000000 27500.000000   NaN   40.395000  2009.000000
75%      1.000000 94250.000000   NaN  132.697500  2011.000000
max      4.000000 730000.000000   NaN  165.000000  2013.000000
```




# Pandas GroupBy: Aggregation

Aggregation returns a reduced version of data

	key	data1	data2
0	A	0	5
1	B	1	0
2	C	2	3
3	A	3	3
4	B	4	7
5	C	5	9

String Function list



```
In [197]: df.groupby('key').aggregate(['min', np.median, max])
```

```
Out[197]:
```

	data1			data2		
	min	median	max	min	median	max
key						
A	0	1.5	3	3	4.0	5
B	1	2.5	4	0	3.5	7
C	2	3.5	5	3	6.0	9

# Pandas GroupBy: Aggregation

Aggregation, Continued:

```
In [198]: df.groupby('key').aggregate({'data1': 'min',  
...:                                'data2': 'max'})
```

```
Out[198]:
```

	data1	data2
key		
A	0	5
B	1	7
C	2	9

# Pandas GroupBy: Filter

Allows us to filter out data that we desire(or drop data we do not desire) based on group properties.

```
In [205]: df
```

```
Out[205]:
```

	key	data1	data2
0	A	0	5
1	B	1	0
2	C	2	3
3	A	3	3
4	B	4	7
5	C	5	9

```
In [206]: df.groupby('key').std()
```

```
Out[206]:
```

	key	data1	data2
A	A	2.12132	1.414214
B	B	2.12132	4.949747
C	C	2.12132	4.242641

Standard deviation greater than 4 for B and C keys.

```
In [207]: def filter_func(x):
```

```
...:     return x['data2'].std() > 4
```

```
...:
```

Filter function: Display data from column 'data2' where standard deviation greater than 4.

```
In [208]: df.groupby('key').filter(filter_func)
```

```
Out[208]:
```

	key	data1	data2
1	B	1	0
2	C	2	3
4	B	4	7
5	C	5	9

Therefore result only includes rows of keys whose standard deviation for column we grouped on – 'key' is greater than 4

# Pandas GroupBy: Transformation

- Transform passes each column for each group as Series to the custom function.
- The custom function passed to transform must return a sequence (a one dimensional Series, array or list) the same length as the group.

```
In [214]: df
```

```
Out[214]:
```

	key	data1	data2
0	A	0	5
1	B	1	0
2	C	2	3
3	A	3	3
4	B	4	7
5	C	5	9

```
In [215]: df.groupby('key').transform(lambda x: x - x)
```

```
Out[215]:
```

	data1	data2
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
5	0	0

```
In [216]: df.groupby('key').transform(lambda x: x - x.mean())
```

```
Out[216]:
```

	data1	data2
0	-1.5	1.0
1	-1.5	-3.5
2	-1.5	-3.0
3	1.5	-1.0
4	1.5	3.5
5	1.5	3.0

# Pandas GroupBy: Apply

Apply implicitly passes all the columns for each group as a **DataFrame** to the custom function – note transform passes each column for each group as a **Series** to the custom function

```
In [227]: df
```

```
Out[227]:
```

	key	data1	data2
0	A	0	5
1	B	1	0
2	C	2	3
3	A	3	3
4	B	4	7
5	C	5	9

```
In [228]: def norm_by_data2(x):
```

```
....:     # x is a DataFrame of group values
....:     x['data1'] /= x['data2'].sum()
....:     return x
....:
```

Normalize column data1 by the sum of data2 column

```
In [229]: df.groupby('key').apply(norm_by_data2)
```

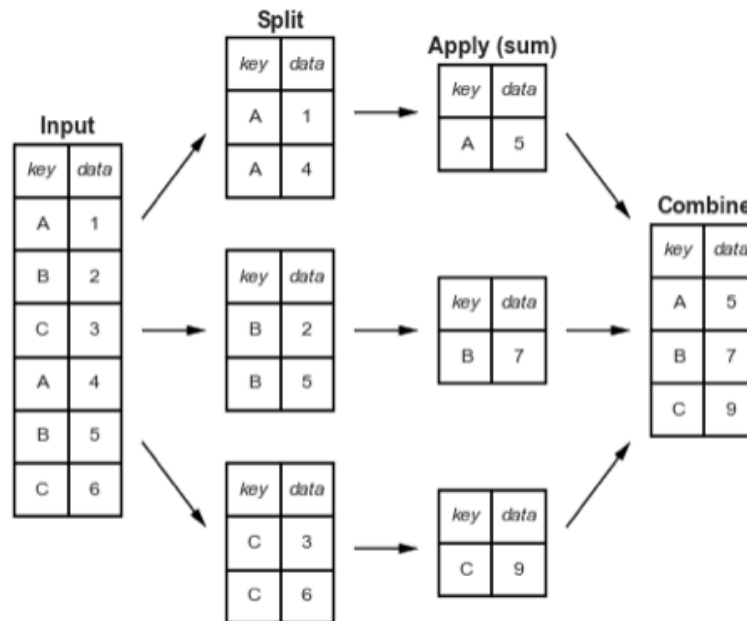
```
Out[229]:
```

	key	data1	data2
0	A	0.000000	5
1	B	0.142857	0
2	C	0.166667	3
3	A	0.375000	3
4	B	0.571429	7
5	C	0.416667	9

Would transform instead of apply work here? Why or why not?

# Pandas: Pivot Tables

Pivot table is a multidimensional version of GroupBy aggregation. Remember split and combine:



In Pivot Table the Split and the Combine are NOT APPLIED ONLY to ONE DIMENSION INDEX – but across two dimensional grid

# Pandas: Pivot Tables

Titanic dataset: Available from seaborn library

- We can apply groupby function to e.g. examine fare(ticket) prices bought by males and females by first grouping on 'sex' column and then on the returned DataFrameGroupBy object calling function describe() on the fare column.

```
In [166]: titanic.head()
Out[166]:
```

	survived	pclass	sex	age	...	deck	embark_town	alive	alone
0	0	3	male	22.0	...	NaN	Southampton	no	False
1	1	1	female	38.0	...	C	Cherbourg	yes	False
2	1	3	female	26.0	...	NaN	Southampton	yes	True
3	1	1	female	35.0	...	C	Southampton	yes	False
4	0	3	male	35.0	...	NaN	Southampton	no	True

[5 rows x 15 columns]

```
In [167]: titanic.groupby('sex')[['fare']].head()
Out[167]:
```

	fare
0	7.2500
1	71.2833
2	7.9250
3	53.1000
4	8.0500
5	8.4583
6	51.8625
7	21.0750
8	11.1333
9	30.0708

```
In [168]: titanic.groupby('sex')[['fare']].describe()
Out[168]:
```

	fare	count	mean	std	min	25%	50%	75%	max
sex									
female		314.0	44.479818	57.997698	6.75	12.071875	23.0	55.00	512.3292
male		577.0	25.523893	43.138263	0.00	7.895800	10.5	26.55	512.3292

# Pandas: Pivot Tables

The information obtained on the prices of the fare tickets paid by male and female passengers was interesting, but what if we would like to look at the same data but now also with information on the fare class (first, second, third)

```
In [170]: titanic.groupby('sex')[['fare']].mean()
```

```
Out[170]:
```

	fare
sex	
female	44.479818
male	25.523893

```
In [171]: titanic.pivot_table('fare', index='sex', columns='class')
```

```
Out[171]:
```

	First	Second	Third
sex			
female	106.125798	21.970121	16.118810
male	67.226127	19.741782	12.661633



# Pandas: TODO

---

- To Do Before Lab2:
  - Read **the following sections** of the Chapter 3:
    - Introduction Pandas Objects
    - Data Indexing and Selection
    - Operating on Data in Pandas
    - Handling Missing Data
    - Aggregation and Grouping
    - Pivot Tables -